

Assignment 3 Solution

Import necessary libraries

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import tensorflow as tf

from keras.models import Sequential
from keras.layers import Convolution2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense

from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator

import cv2
```

Image Augmentation

In [2]:

```
data_path = './flowers/'
batch_size = 32
target_size = (64, 64)
```

In [3]:

```
train_datagen = ImageDataGenerator(rescale=1./255,
                                   shear_range=0.2,
                                   zoom_range=0.2,
                                   width_shift_range=0.1,
                                   height_shift_range=0.1,
                                   horizontal_flip=True,
                                   validation_split=0.2)

test_datagen = ImageDataGenerator(rescale=1. / 255, validation_split=0.2)
```

In [4]:

```
X_train = train_datagen.flow_from_directory(data_path,
                                             target_size=target_size,
                                             batch_size=batch_size,
                                             subset="training",
                                             class_mode='categorical')

X_test = test_datagen.flow_from_directory(data_path,
                                           target_size=target_size,
                                           batch_size=batch_size,
                                           subset="validation",
                                           class_mode='categorical')
```

Found 3457 images belonging to 5 classes.

Found 860 images belonging to 5 classes.

Create model

In [6]:

```
model = Sequential()
```

Add layers

In [7]:

```
model.add(Convolution2D(32, (3, 3), input_shape=(64, 64, 3),  
activation='relu'))  
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Convolution2D(32, (3, 3), activation='relu'))  
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Convolution2D(64, (3, 3), activation='relu'))  
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Flatten())
```

```
model.add(Dense(units=64, activation='relu'))  
model.add(Dense(units=5, activation='softmax'))
```

In [8]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 29, 29, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 64)	147520
dense_1 (Dense)	(None, 5)	325
=====		
Total params: 176,485		
Trainable params: 176,485		
Non-trainable params: 0		

Compile the model

In [9]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy',  
metrics=['accuracy'])
```

Fit the model

In [12]:

```
model.fit(X_train, steps_per_epoch=100, epochs=20)  
Epoch 1/20  
100/100 [=====] - 8s 74ms/step - loss: 0.8592 - ac  
curacy: 0.6693  
Epoch 2/20  
100/100 [=====] - 8s 75ms/step - loss: 0.8127 - ac  
curacy: 0.6841  
Epoch 3/20  
100/100 [=====] - 7s 75ms/step - loss: 0.7724 - ac  
curacy: 0.7009  
Epoch 4/20  
100/100 [=====] - 7s 73ms/step - loss: 0.7421 - ac  
curacy: 0.7198  
Epoch 5/20  
100/100 [=====] - 8s 78ms/step - loss: 0.7266 - ac  
curacy: 0.7204  
Epoch 6/20  
100/100 [=====] - 8s 75ms/step - loss: 0.7393 - ac  
curacy: 0.7160  
Epoch 7/20  
100/100 [=====] - 8s 82ms/step - loss: 0.7395 - ac  
curacy: 0.7163  
Epoch 8/20  
100/100 [=====] - 8s 80ms/step - loss: 0.6993 - ac  
curacy: 0.7251  
Epoch 9/20  
100/100 [=====] - 8s 76ms/step - loss: 0.6899 - ac  
curacy: 0.7444  
Epoch 10/20  
100/100 [=====] - 8s 76ms/step - loss: 0.6781 - ac  
curacy: 0.7457  
Epoch 11/20  
100/100 [=====] - 8s 76ms/step - loss: 0.6612 - ac  
curacy: 0.7463  
Epoch 12/20  
100/100 [=====] - 8s 76ms/step - loss: 0.6526 - ac  
curacy: 0.7539  
Epoch 13/20  
100/100 [=====] - 8s 76ms/step - loss: 0.6387 - ac  
curacy: 0.7485  
Epoch 14/20  
100/100 [=====] - 8s 77ms/step - loss: 0.6011 - ac  
curacy: 0.7712  
Epoch 15/20  
100/100 [=====] - 8s 76ms/step - loss: 0.5926 - ac  
curacy: 0.7826  
Epoch 16/20
```

```

100/100 [=====] - 8s 80ms/step - loss: 0.6164 - ac
curacy: 0.7608
Epoch 17/20
100/100 [=====] - 8s 77ms/step - loss: 0.6137 - ac
curacy: 0.7690
Epoch 18/20
100/100 [=====] - 8s 75ms/step - loss: 0.5976 - ac
curacy: 0.7703
Epoch 19/20
100/100 [=====] - 8s 80ms/step - loss: 0.6437 - ac
curacy: 0.7545
Epoch 20/20
100/100 [=====] - 8s 76ms/step - loss: 0.6004 - ac
curacy: 0.7700

```

Out[12]:

```
<keras.callbacks.History at 0x274702feb90>
```

Save the model

In [13]:

```
model.save("model.h5")
```

Test the model

In [17]:

```

def predict():
    img = image.load_img("./rose.jpg", target_size=target_size)
    x = image.img_to_array(img)
    x = tf.expand_dims(x,0)

    labels = ['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']

    pred = model.predict(x)
    prediction = labels[np.argmax(pred[0])]

    print(f'The given image is a {prediction}')
    plt.imshow(plt.imread("./rose.jpg"))
    plt.axis('off')
    plt.show()

```

In [18]:

```

predict()

1/1 [=====] - 0s 22ms/step
The given image is a rose

```

